

## Sparsity in images, a basis better than Fourier basis

### Background knowledge: Cooley-Tukey FFT

The *Fourier basis* in  $\mathbb{R}^n$  is given by the DFT matrix  $(F_n)_{i,j} = \omega^{(i-1)(j-1)}$ , where  $\omega = e^{-2\pi i/n}$  is the primitive  $n$ -th root of unity.  $F_n$  is a Vandermonde matrix:

$$F_n = \begin{pmatrix} 1 & 1 & 1 & \dots & 1 & 1 & \dots & 1 & 1 \\ 1 & \omega & \omega^2 & \dots & \omega^{\frac{n}{2}-1} & \omega^{\frac{n}{2}} & \dots & \omega^{n-2} & \omega^{n-1} \\ 1 & \omega^2 & \omega^4 & \dots & \omega^{n-2} & \omega^n & \dots & \omega^{2n-4} & \omega^{2n-2} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 1 & \omega^{\frac{n}{2}-1} & \omega^{n-2} & \dots & \omega^{(\frac{n}{2}-1)^2} & \omega^{(\frac{n}{2}-1)\frac{n}{2}} & \dots & \omega^{(\frac{n}{2}-1)(n-2)} & \omega^{(\frac{n}{2}-1)(n-1)} \\ 1 & \omega^{\frac{n}{2}} & \omega^n & \dots & \omega^{\frac{n}{2}(\frac{n}{2}-1)} & \omega^{(\frac{n}{2})^2} & \dots & \omega^{\frac{n}{2}(n-2)} & \omega^{\frac{n}{2}(n-1)} \\ \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 1 & \omega^{n-1} & \omega^{2n-2} & \dots & \omega^{(n-1)(\frac{n}{2}-1)} & \omega^{(n-1)\frac{n}{2}} & \dots & \omega^{(n-1)(n-2)} & \omega^{(n-1)^2} \end{pmatrix}. \quad (1)$$

The Cooley-Tukey FFT is a divide-and-conquer algorithm for computing the DFT. For simplicity, we assume  $n$  is a power of 2, such that we can divide the matrix into 4 blocks:

- The odd columns (blue background), top half:

$$F_{\text{odd, top}} = \begin{pmatrix} 1 & 1 & \dots & 1 \\ 1 & \omega^2 & \dots & \omega^{n-2} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & \omega^{n-2} & \dots & \omega^{(\frac{n}{2}-1)(n-2)} \end{pmatrix} = \begin{pmatrix} 1 & 1 & \dots & 1 \\ 1 & (\omega^2) & \dots & (\omega^2)^{\frac{n}{2}-1} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & (\omega^2)^{\frac{n}{2}-1} & \dots & (\omega^2)^{(\frac{n}{2}-1)(\frac{n}{2}-1)} \end{pmatrix} = F_{\frac{n}{2}} \quad (2)$$

- The even columns (white background), top half:

$$F_{\text{even, top}} = D_{\frac{n}{2}} F_{\frac{n}{2}} \quad (3)$$

where  $D_n = \text{diag}(1, \omega, \omega^2, \dots, \omega^{n-1})$ .

- The odd columns (blue background), bottom half:

$$F_{\text{odd, bottom}} = F_{\frac{n}{2}}. \quad (4)$$

Note  $\omega^n = 1$  is ignored.

- The even columns (white background), bottom half:

$$F_{\text{even, bottom}} = -D_{\frac{n}{2}} F_{\frac{n}{2}}, \quad (5)$$

where the minus sign comes from  $\omega^{\frac{n}{2}} = -1$ .

Finally, we arrive at the Cooley-Tukey FFT given by:

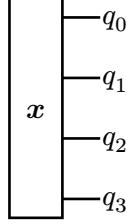
$$F_n \mathbf{x} = \begin{pmatrix} I_{\frac{n}{2}} & D_{\frac{n}{2}} \\ I_{\frac{n}{2}} & -D_{\frac{n}{2}} \end{pmatrix} \begin{pmatrix} F_{\frac{n}{2}} & 0 \\ 0 & F_{\frac{n}{2}} \end{pmatrix} \begin{pmatrix} \mathbf{x}_{\text{odd}} \\ \mathbf{x}_{\text{even}} \end{pmatrix} \quad (6)$$

where  $\mathbf{x}_{\text{odd}}$  and  $\mathbf{x}_{\text{even}}$  contain the odd and even indexed elements of  $\mathbf{x}$ , respectively. It indicates that the discrete Fourier transformation in  $\mathbb{R}^n$  can be decomposed into two smaller discrete Fourier transformations in  $\mathbb{R}^{\frac{n}{2}}$  with a diagonal matrix  $D_n$  in between. Note applying diagonal matrices can be done in  $O(n)$  operations, this decomposition leads to the recurrence relation  $T(n) = 2T(\frac{n}{2}) + O(n)$ , which solves to  $O(n \log n)$  total operations.

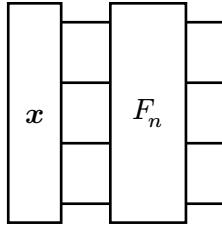
The inverse transformation is given by  $F_n^\dagger \mathbf{x}/n$ . The DFT matrix is unitary up to a scale factor:  $F_n F_n^\dagger = nI$ .

## Tensor network representation of the Cooley-Tukey FFT

This section requires preliminary knowledge of tensor networks (TODO: add reference). In tensor network diagram, a vector of size  $n = 2^k$  can be represented as a tensor with  $k$  indices, denoting the basis index  $i = 2^0 q_0 + 2^1 q_1 + \dots + 2^{k-1} q_{k-1}$ .

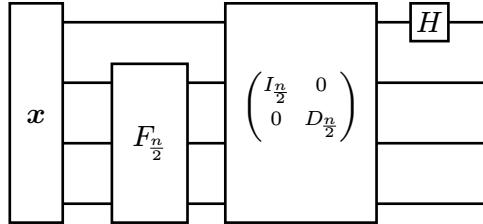


In the following, we aim to find a tensor network decomposition for the linear map  $F_n$ :



Step 1: To start, the equation Equation 6 can be represented as the following tensor network:

$$F_n x = \left( \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix} \otimes I_{\frac{n}{2}} \right) \begin{pmatrix} I_{\frac{n}{2}} & 0 \\ 0 & D_{\frac{n}{2}} \end{pmatrix} \begin{pmatrix} F_{\frac{n}{2}} x_{\text{odd}} \\ F_{\frac{n}{2}} x_{\text{even}} \end{pmatrix}, \quad (7)$$

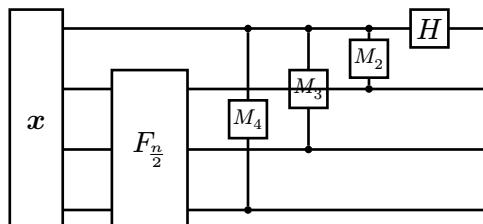


where  $H = \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$  is a Hadamard matrix (upto a constant factor).

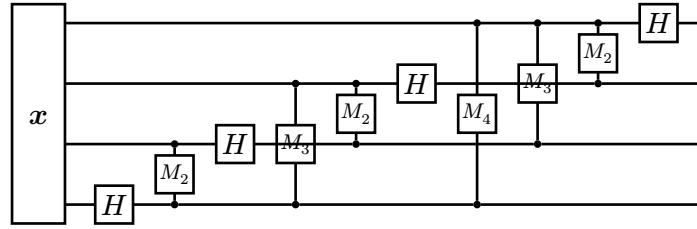
Step 2: Then, we will decompose the diagonal matrix  $\begin{pmatrix} I_{\frac{n}{2}} & 0 \\ 0 & D_{\frac{n}{2}} \end{pmatrix}$  into a tensor network. This diagonal matrix corresponds to operation: if  $q_0$  is 0 (odd index), the no operation is applied, otherwise (even index) the operation  $D_{\frac{n}{2}}$  is applied. Observe that  $D_n = \text{diag}(1, \omega^{\frac{n}{2}}) \otimes \text{diag}(1, \omega, \omega^2, \dots, \omega^{\frac{n}{2}-1}) = \text{diag}(1, \omega^{\frac{n}{2}}) \otimes \text{diag}(1, \omega^{\frac{n}{4}}) \otimes \dots \otimes \text{diag}(1, \omega)$ . We have

$$\begin{pmatrix} I_{\frac{n}{2}} & 0 \\ 0 & D_{\frac{n}{2}} \end{pmatrix} = \text{ctrl}_0 \left( \text{diag}(1, \omega^{\frac{n}{4}})_1 \right) \text{ctrl}_0 \left( \text{diag}(1, \omega^{\frac{n}{8}})_2 \right) \dots \text{ctrl}_0 \left( \text{diag}(1, \omega)_{\log_2 n} \right), \quad (8)$$

where  $\text{ctrl}_i(A_j)$  means the target operation applied on  $A_j$  is applied only if bit  $q_i$  is 1. Here, since the controlled gate is diagonal, it can be represented as a matrix connecting two variables:



In this diagram,  $M_k = \begin{pmatrix} 1 & 1 \\ 1 & e^{i\pi/2^{k-1}} \end{pmatrix}$  connects the two qubits involved in the controlled operation, which effectively multiplies a phase factor  $e^{i\pi/2^{k-1}}$  if two bit are both in state 1. By recursively decomposing the  $F_{\frac{n}{2}}$  tensor, we can obtain the following tensor network.

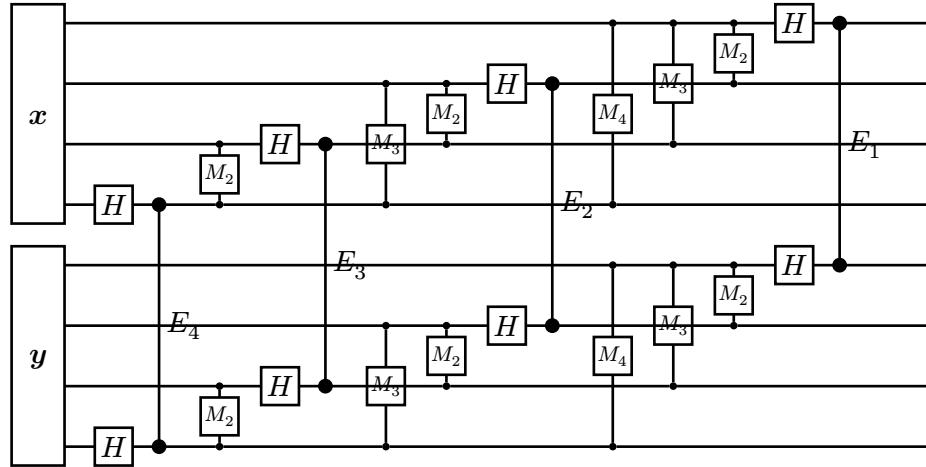


Direct evaluation of this tensor network takes  $O(n \log^2 n)$  operations. By respecting the fact that the *controlled phase* operation is a diagonal matrix, we can merge these operations and further reduce the complexity to  $O(n \log(n))$ .

### Entangled Fourier Basis: XY Correlation

In the standard 2D Fourier transform, the  $x$  and  $y$  coordinates are processed independently. For an image of size  $2^n \times 2^n$  (i.e., square images with  $m = n$ ), we apply QFT on the  $n$  row qubits and separately on the  $n$  column qubits. This independence assumption is often suboptimal for natural images where spatial correlations exist between rows and columns.

We propose an *entangled QFT basis* that introduces controlled-phase gates between  $x$  and  $y$  qubits after each layer of the QFT circuit. For the square case  $m = n$ , we use a *one-to-one* entanglement structure where each  $x$  qubit  $x_k$  is coupled with the corresponding  $y$  qubit  $y_k$ . This creates correlation between the two spatial dimensions:



The entanglement gates  $E_k = \text{diag}(1, 1, 1, e^{i\varphi_k})$  are parameterized controlled-phase gates that couple the  $k$ -th qubit from the  $x$ -axis with the  $k$ -th qubit from the  $y$ -axis. For a square  $2^n \times 2^n$  image encoded with  $n$  row qubits and  $n$  column qubits (total  $2n$  qubits) and one-to-one coupling between corresponding row/column qubits, we add exactly  $n$  entanglement gates, one after each Hadamard layer.

The total transformation becomes:

$$\mathcal{T}_{\text{entangled}} = U_{\text{entangle}} \cdot (F_n \otimes F_n) \quad (9)$$

where  $U_{\text{entangle}} = \prod_{k=1}^n E_k$  is the product of all entanglement gates, and  $F_n$  is the  $n$ -qubit QFT applied along each spatial dimension.

Key advantages of this approach:

- Captures diagonal features and cross-dimensional patterns common in natural images
- Maintains  $O(n \log n)$  computational complexity in the number  $n$  of qubits per spatial dimension (equivalently  $O(N \log N)$  for linear image size  $N = 2^n$ ), matching the standard QFT
- Adds exactly  $n$  additional real-valued learnable parameters  $\varphi_k$  (one phase per qubit pair), i.e.,  $O(n)$  in  $n$
- Reduces to standard 2D QFT when all entanglement phases  $\varphi_k = 0$

### Alternative Basis: Time Evolving Block Decimation (TEBD)

Time Evolving Block Decimation (TEBD) is a tensor network ansatz originally developed for simulating 1D quantum many-body systems. It employs a *brickwork* pattern of nearest-neighbor two-qubit gates applied in alternating layers. For image processing on a  $2^n \times 2^n$  grid, TEBD can be adapted by treating the  $n$  row qubits and  $n$  column qubits as two coupled 1D chains.

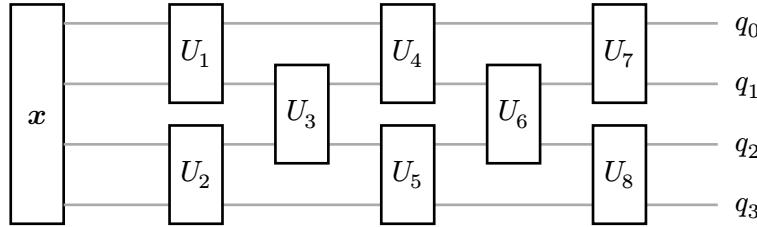


Figure 8: TEBD brickwork circuit for  $n = 4$  qubits with  $L = 5$  layers: alternating layers of nearest-neighbor two-qubit gates  $U_k$ .

In this diagram, each  $U_k$  is a parameterized  $4 \times 4$  unitary matrix acting on two adjacent qubits. Common parameterization choices for  $U_k$  include:

1. **Full  $U(4)$  unitary:** The most general form with 16 complex parameters constrained by unitarity ( $UU^\dagger = I$ ). This lies on the unitary manifold  $U(4)$ .
2. **Hardware-efficient ansatz:** Decompose each two-qubit gate as single-qubit rotations followed by an entangling gate:

$$U_k = (R_z(\varphi_1)R_y(\theta_1) \otimes R_z(\varphi_2)R_y(\theta_2)) \cdot \text{CZ} \cdot (R_z(\varphi_3)R_y(\theta_3) \otimes R_z(\varphi_4)R_y(\theta_4)) \quad (10)$$

where  $R_y(\theta) = \exp(-i\theta \frac{Y}{2})$ ,  $R_z(\varphi) = \exp(-i\varphi \frac{Z}{2})$ , and CZ is the controlled-Z gate. This uses 8 real parameters per gate.

3. **XX+YY+ZZ interaction:** Inspired by Hamiltonian simulation:

$$U_k = \exp(i(\alpha_k X \otimes X + \beta_k Y \otimes Y + \gamma_k Z \otimes Z)) \quad (11)$$

with only 3 real parameters  $(\alpha_k, \beta_k, \gamma_k)$  controlling the entanglement strength.

Direct evaluation of this tensor network takes  $O(nL \cdot 4^2) = O(nL)$  operations for  $L$  layers. The parameter space consists of  $\lfloor L/2 \rfloor \lfloor n/2 \rfloor + \lceil L/2 \rceil \lfloor (n-1)/2 \rfloor$  two-qubit gates (approximately  $(n/2) \cdot L$  gates for large  $n$ ). The total parameter manifold is:

$$\mathcal{M}_{\text{TEBD}} = \prod_{k=1}^{\text{|gates|}} U(4) \quad (12)$$

For Riemannian optimization, we optimize on this product of unitary manifolds using the same gradient descent approach as the QFT basis.

## Alternative Basis: Multi-scale Entanglement Renormalization Ansatz (MERA)

The Multi-scale Entanglement Renormalization Ansatz (MERA) is a hierarchical tensor network that naturally captures *multi-scale correlations*. It consists of alternating layers of *disentanglers* (two-qubit unitaries) and *isometries* (coarse-graining maps), forming a tree-like structure. For  $n = 2^k$  qubits, MERA has  $k$  layers, with each layer reducing the number of effective qubits by half.

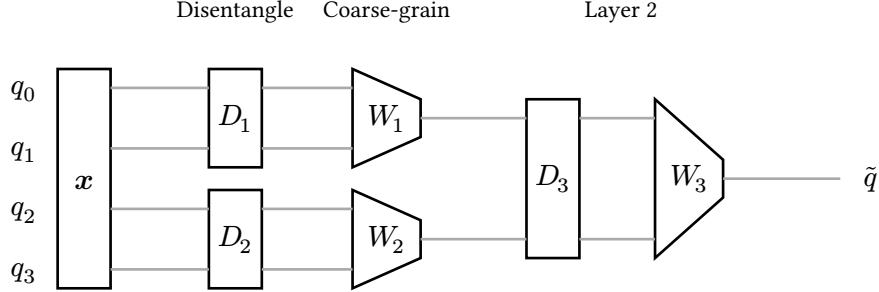


Figure 9: MERA circuit for  $n = 4$  qubits: disentanglers  $D_k$  (rectangles) remove short-range entanglement, isometries  $W_k$  (trapezoids) perform 2-to-1 coarse-graining. Each layer halves the number of effective qubits.

In this diagram, there are two types of parameterized gates:

1. **Disentanglers**  $D_k \in U(4)$ : Parameterized  $4 \times 4$  unitary matrices acting on two adjacent qubits before coarse-graining. Same parameterization options as TEBD gates (full  $U(4)$ , hardware-efficient, or  $XX+YY+ZZ$ ).
2. **Isometries**  $W_k : \mathbb{C}^4 \rightarrow \mathbb{C}^2$ : Parameterized  $2 \times 4$  matrices that map two qubits to one qubit (coarse-graining). They satisfy the isometry constraint  $WW^\dagger = I_2$ , lying on the Stiefel manifold  $\text{St}(2, 4)$ . Parameterization options include:
  - **Full Stiefel**: Any  $2 \times 4$  matrix with orthonormal rows, 8 real parameters
  - **Structured**:  $W = \begin{pmatrix} \cos \theta & \sin \theta e^{i\varphi_1} & 0 & 0 \\ 0 & 0 & \cos \psi & \sin \psi e^{i\varphi_2} \end{pmatrix}$  with 4 real parameters (block-diagonal structure)

Direct evaluation of this tensor network takes  $O(n)$  operations total, since each layer processes  $O(2^{k-l})$  qubits at level  $l$  and there are  $k = \log_2 n$  layers. For  $n = 2^k$  input qubits, the parameter count is:

- Layer  $l$ :  $2^{k-l-1}$  disentanglers +  $2^{k-l-1}$  isometries
- Total:  $\sum_{l=0}^{k-1} 2^{k-l-1} = n - 1$  disentanglers and  $n - 1$  isometries

The total parameter manifold is:

$$\mathcal{M}_{\text{MERA}} = \left( \prod_{k=1}^{n-1} U(4) \right) \times \left( \prod_{k=1}^{n-1} \text{St}(2, 4) \right) \quad (13)$$

For Riemannian optimization, we optimize on this product of unitary and Stiefel manifolds. The hierarchical structure naturally captures multi-scale features (edges → textures → objects), with only  $O(\log n)$  depth for  $n$  qubits.

### Learning a better Fourier basis

Observing that in this representation, tensor parameters can be tuned without affecting the computational complexity, e.g. the parameters in  $M_k$  and  $H$ . Can we find a transformation better than the Fourier basis? Or is Fourier basis already optimal for image processing?

Intuitively, the fourier basis is not optimal for image processing, because:

- the fourier basis assumes periodic boundary condition, which is not suitable for image processing.
- the 2d fourier basis assumes the  $X$  and  $Y$  coordinates are independent, which is not suitable for image processing.

## Tasks

- Create an image dataset  $\mathcal{D} = \{\mathbf{x}_i\}_{i=1}^N$ .
- Create a tensor network transformation based on the above QFT circuit, denoted as  $\mathcal{T}(\boldsymbol{\theta})$ , where  $\boldsymbol{\theta}$  is the parameters of the tensor network.
- Variationally optimize the circuit parameters to capture the sparsity of the image. The cost function is

$$\mathcal{L}(\boldsymbol{\theta}) = \sum_{i=1}^N \|\mathbf{x}_i - \mathcal{T}(\boldsymbol{\theta})^{-1}(\text{truncate}(\mathcal{T}(\boldsymbol{\theta})(\mathbf{x}_i), k))\|_2^2 \quad (14)$$

Here, we can choose a different loss function to capture details in the image, e.g. the edges. For simplicity, we use the  $l_1$ -norm instead:

$$\mathcal{L}(\boldsymbol{\theta}) = \sum_{i=1}^N \|\mathcal{T}(\boldsymbol{\theta})(\mathbf{x}_i)\|_1 \quad (15)$$

This loss will encourage the tensor network to output a sparse pattern in the “moment space”. It is a standard trick that widely used in *compressed sensing*.

- In the 2D Fourier transformation, the  $X$  and  $Y$  coordinates are independent. Here we allow  $X$  and  $Y$  coordinates to correlate with each other in the tensor network basis.
- Add edge detection features.
- Compare the performance with the Fourier basis.